STUDLYF — Feature Specification for Development

Developer reference document — full feature set, explained

How to read this document

This covers every feature in the original Studlyf vision, explained in plain language, with what the founder inputs, what the system needs to output, a suggested technical approach, and at least one reference website to study for inspiration. It's meant to sit alongside the rescoped lean v1 PRD shared earlier — that document defines build order (Funding Navigator, Roadmap, and AI Copilot ship first); this document is the full reference so the whole team understands where everything else fits and why, before it gets built.

A note on the reference sites throughout: they're there to study UI patterns, data structure, and positioning — not to scrape or copy. Several (Crunchbase, Tracxn, OpenVC) hold their own proprietary datasets; Studlyf's data needs to be sourced independently, as detailed feature by feature below.

---

0. Onboarding: Account & Startup Profile Creation

What it does: Gets a founder from "never heard of Studlyf" to "has a generated dashboard" in two short steps.

Step 1 — Create Account: Google login or email login, with optional LinkedIn import or resume upload to pre-fill profile fields.

Step 2 — Create Startup Profile: Founder enters startup name, description, industry, country, and selects a stage (Idea, Validation, MVP, Revenue, Fundraising).

Feature 1 — AI Startup Validation

What it does: Helps a founder figure out whether their idea is actually worth pursuing, before they spend months building it.

Founder inputs: Startup idea, problem statement, customer segment, geography.

System outputs:

- Market research: market size, growth trends, industry overview
- Competitor discovery: direct and indirect competitors, funding history, pricing
- Customer persona: ideal customer profile, pain points, behavior
- Validation score (0–100), broken into demand, competition, scalability, and revenue potential sub-scores

Suggested approach: Market research and competitor discovery need to be current, so this should be built on an LLM with live web search/retrieval grounding rather than a static internal database — markets move too fast for a one-time-built dataset to stay accurate. The validation score should always be shown with its four sub-scores visible, never just the single 0–100 number; a single number with no visible basis is the kind of thing sophisticated founders (and investors looking over their shoulder) will distrust immediately.

Reference sites: Crunchbase for how competitor and funding-history data is structured and displayed.

Feature 2 — Founder Roadmap ("Founder GPS")

What it does: Tells founders exactly what to do next, week by week, based on their current stage. This is the most important feature in the whole product — it's the literal "GPS" in the positioning.

Founder input: Current stage (e.g., Idea Stage).

System output: A week-by-week task list (e.g., Week 1: conduct 20 interviews; Week 2: build a landing page; Week 3: collect feedback; Week 4: create MVP), with each task linked to templates, guides, and resources.

Suggested approach: Build this as a curated content library first — a structured set of stage-specific playbooks written and vetted by the team, stored as structured data (a CMS or simple database table), not generated fresh by an AI each time. Vetted, evergreen content is more trustworthy than freshly-generated text for something this central, and it's far easier to keep consistent. AI can be layered in later to personalize sequencing or ordering based on the founder's sector and signals, once the base content exists.

Reference sites: Founder Institute for the structured curriculum/playbook model; Notion for the task-and-template UI pattern (checklists, linked documents, progress tracking).

Feature 3 — AI Build Advisor

What it does: Helps founders figure out what to actually build — tech stack, build phases, and rough costs.

Founder input: Startup type (SaaS, AI Startup, Marketplace, Mobile App).

System output: Recommended stack (frontend, backend, database, hosting, AI provider), a phased development roadmap (MVP, Growth Features, Scale Features), and a cost estimator (development, hosting, AI costs).

Suggested approach: This can be a straightforward rules-based recommendation engine (a decision tree keyed on startup type) rather than something needing fresh AI generation each time — stack recommendations don't need to be novel. The cost estimator should pull from real, current pricing (actual hosting and API pricing pages) rather than static guessed figures, since those numbers change.

Reference sites: Lovable, v0, and Bolt. Worth being honest internally about this one: these three already do AI-assisted app generation extremely well, and Lovable in particular has become a major, well-funded player. A static stack-recommendation feature next to those is thin. Consider scoping this as a lightweight advisor that helps a non-technical founder understand which of those tools fits their situation and why, rather than trying to rebuild what they already do.

Feature 4 — Funding Navigator

What it does: Helps founders discover the funding opportunities that actually fit them — incubators, accelerators, government programs, grants, fellowships.

Founder inputs: Sector, stage, geography.

System outputs: Relevant matches across each category, an eligibility indicator, and an application tracker (Applied / Under Review / Accepted / Rejected).

Suggested approach: This needs to be a manually curated, verified database — structured fields per program (stage range, sector tags, geography, deadline, equity terms, application link, last-verified date) — not a scrape of a third-party site. Show eligibility as a transparent breakdown (which specific criteria the founder meets, which they don't) rather than a single fabricated percentage like "90% eligible," which implies a level of confidence the system can't actually have without real historical acceptance data. The application tracker is a simple status field tied to the founder's account, updated by the founder themselves initially.

Reference sites: Startup India for how scheme structure and eligibility criteria are typically organized; OpenVC for how a funding-discovery-plus-tracking interface can be laid out (note: OpenVC already publishes a large free database of investors and accelerators — study the UX, but Studlyf's actual data still needs to be built independently and verified, especially for India-specific government schemes that a global tool won't cover well).

Feature 5 — Investor Discovery

What it does: Helps founders identify investors that actually fit their industry, stage, and revenue profile.

Founder inputs: Industry, stage, revenue.

System outputs: Investor name, stage preference, geography, focus areas; a fundraising readiness score across team, product, market, and traction.

Suggested approach: Be aware going in that this category already has strong, well-resourced incumbents with large existing datasets. Rather than building a proprietary investor database from scratch immediately, consider a partnership or licensed-data approach for the initial version, and reserve in-house data-building effort for the areas (like India-specific grants) where the market is genuinely underserved.

Reference sites: Tracxn, Crunchbase, and OpenVC — all three already do this well; study what they don't cover (especially India-specific early-stage angel and micro-VC data) for a possible wedge.

Feature 6 — Mentor Discovery

What it does: Helps founders find mentors relevant to their industry, stage, and geography.

Filter criteria: Industry, startup stage, geography.

System output: Mentor profile showing expertise, startup experience, and availability.

Suggested approach: Start as a curated, manually-onboarded directory rather than an automated matching algorithm — match quality here depends entirely on how deep and accurate the mentor data is, which needs human curation, especially early on. Mentors are also a natural thing to source through incubator and accelerator partnerships rather than building cold.

Reference sites: Founder Institute for its mentor program structure (expertise tagging, availability, matching model).

Feature 7 — Relationship Intelligence

What it does: Identifies people a founder should talk to next — potential customers, investors, mentors, other founders — and shows a path to a warm introduction through mutual contacts.

Founder inputs: Resume, contact list, LinkedIn export.

System output: A list of relevant people to connect with, and a visualized introduction path (Founder → Mutual Contact → Investor).

Important build note before any code is written: Do not build this on scraped LinkedIn data. LinkedIn's user agreement explicitly prohibits automated scraping of profile data, and enforcement has been active and escalating — including litigation where the requested remedy covered not just the scraped data itself but anything inferred, aggregated, or synthesized from it. That means even a derived "introduction path" graph built on top of scraped profiles carries real legal exposure, not just the raw scrape. Build this instead on data the founder explicitly owns and uploads with consent (their own exported contacts, a CSV they control) and on mutual-introduction requests where both sides opt in — never on inferring a path through someone else's network without their knowledge. Get legal review (India's DPDP Act and any other relevant jurisdiction) before development starts on this feature specifically.

Reference sites: Clay and Common Room for relationship-mapping UI patterns — but study their compliant, consent-based data-sourcing approach as closely as their interface, since that's the part that actually matters here.

Feature 8 — Opportunity Radar

What it does: Keeps founders passively informed about what's happening in the ecosystem — trending sectors, new grants, new accelerators and incubators, innovation challenges, startup competitions.

System output: A feed of trending sectors (e.g., AI, Healthcare, Defence, Manufacturing, Climate Tech), new funding programs, and competitions.

Suggested approach: This is largely a byproduct of Feature 4 once that database exists — new or updated entries in the Funding Navigator dataset naturally become radar items. Layer in a lightweight editorial/news component for sector trends and competitions on top of that same data store, rather than building a separate system.

Reference sites: Product Hunt for feed and launch-card UI patterns; YourStory for India-specific startup news tone and structure.

Feature 9 — AI Founder Copilot

What it does: A conversational AI startup advisor founders can ask anything — "How do I get my first 100 users?", "Which incubator should I apply to?" — positioned as the homepage search experience.

Suggested approach: Build this as an LLM-based chat interface grounded primarily in Studlyf's own verified internal data (Funding Navigator entries, Roadmap content, Resources Library) through retrieval, so that specific factual questions — "Am I eligible for X grant?" — pull from verified, sourced answers rather than the model's general knowledge. Broader advice questions (pricing strategy, growth tactics) can draw on the model's general knowledge, but should be visibly distinguished from sourced, database-backed answers so founders know which kind of answer they're getting.

Suggested approach for reference UX: look at any modern retrieval-grounded AI chat product for the "answer plus visible source" display pattern — the key UX principle is that the founder should be able to tell, at a glance, whether an answer comes from Studlyf's verified database or from general AI knowledge.

Build note: Per the lean v1 PRD, this should function as the actual front door of the product — the homepage — not a deep separate feature buried in navigation.

Feature 10 — Startup Resources Library

What it does: A library of practical templates and guides — legal templates, pitch deck templates, founder guides, grant application templates, financial models.

Suggested approach: A structured, versioned content library (a simple CMS) that can launch with a small, genuinely useful curated set rather than trying to be exhaustive on day one. Quality and accuracy matter more than volume here.

Reference sites: Y Combinator's Startup Library for content organization, tone, and depth.

Build note: Anything India-specific in the legal template category (incorporation documents, cap table templates, employment agreements) should be reviewed by an actual lawyer before publishing. A generic AI-generated legal template that a founder relies on and that turns out to be wrong is a real liability, not just a quality issue.

Feature 11 — Founder Profile

What it does: A founder's personal record — startup progress, validation reports, funding applications, roadmap progress, and AI conversation history.

Suggested approach: This is purely an aggregation and presentation layer over data already generated by the other features — no new data source is needed. It's essentially an account dashboard.

Reference sites: No external product needed for inspiration here beyond a clean SaaS account/settings page pattern; this is the lowest-risk, lowest-priority feature to build, since it has no dependency on external data.

Feature 12 — Startup Profile

What it does: A startup-level record — team details, industry, stage, validation reports, funding status.

Suggested approach: Same as Founder Profile — an internal aggregation layer. If the business model evolves toward partnering with incubators (selling them a view into applicant startup profiles), this page could later become the basis for a partner-facing dashboard — worth keeping that possibility in mind in the data model, even if the UI stays private-only at first.

Reference sites: Crunchbase company profile pages, as a layout reference for how a structured startup profile can be organized.

Build note: Keep this private by default. Only make any part of it public or partner-visible for a clear, specific reason, since startup and founder data carries its own privacy considerations.

Admin Panel

What it does: Lets the internal team manage founders, startups, investors, grants, incubators, mentors, AI prompts, and the opportunity feed.

Suggested approach: This is purely operational tooling, not user-facing — there's little reason to build it bespoke from scratch. An off-the-shelf internal admin tool builder (e.g., Retool) or a standard framework admin panel can get this running fast, leaving engineering time for the user-facing product.

Reference sites: Retool as an example of how internal admin tooling is typically assembled quickly.

Build note: This is the tool the ops team will use daily to keep the Funding Navigator data fresh and accurate — it's worth prioritizing right alongside Feature 4, not after it, since the data-quality plan depends on it.

Navigation & Information Architecture

The product is constrained to five primary navigation items: Dashboard, Roadmap, Funding, Network, AI Copilot. Here's a suggested mapping of the twelve features into that structure:

- **Dashboard** → AI Founder Dashboard (Startup Health Score, Validation Status, Recommended Actions), Founder Profile (Feature 11), Startup Profile (Feature 12)
- **Roadmap** → Founder Roadmap (Feature 2), with Resources Library (Feature 10) surfaced contextually inside roadmap tasks as linked templates and guides
- **Funding** → Funding Navigator (Feature 4), with Opportunity Radar (Feature 8) as a feed within the same section
- **Network** → Investor Discovery (Feature 5), Mentor Discovery (Feature 6), and Relationship Intelligence (Feature 7), as tabs within one section

- **AI Copilot** → Feature 9, also reachable as a persistent search/chat entry point from every other screen

AI Startup Validation (Feature 1) and AI Build Advisor (Feature 3) don't need their own nav slots — they fit naturally as actions launched from the Dashboard or Roadmap (e.g., a "Validate my idea" or "Get a build plan" action card), keeping the five-item navigation rule intact.